

* Data Mining:

① Discovering patterns, insights and knowledge from data.

② Process large amount of data quickly.

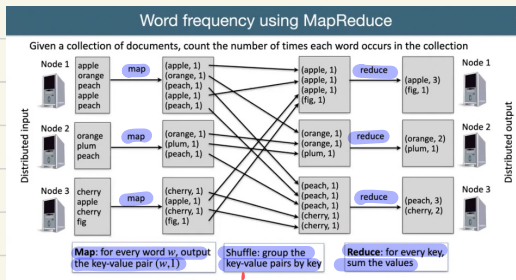

→ The data to be processed is very large and hence stored onto different locations.

Traditional Parallelism

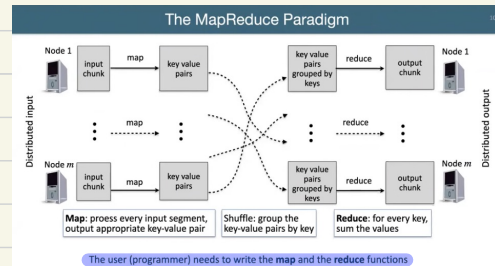
° Bring data to the computer

MapReduce

° Compute where data is present



↓
GroupBy will be done by the system (No need to code)



Note: HDFS runs on top of operating system.
eg.: Implementing HDFS on Linux.

* Hadoop: Open Source MapReduce platform
2 components:

- HDFS: Hadoop Distributed File System (used to store data)
- MapReduce engine: (to process data)

It uses Master-Slave architecture using commodity servers.

HDFS

MapReduce



Master: Namenode (Metadata)

Master: JobTracker

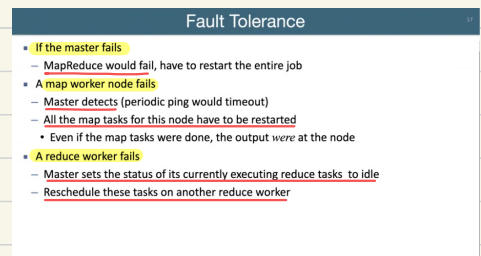
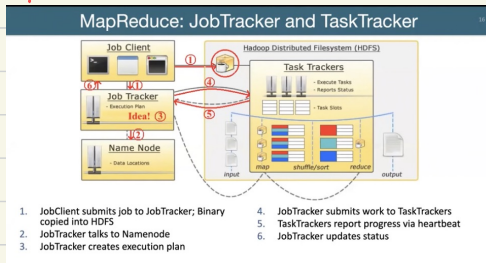
Slave: Datanode (Data)

Slave: TaskTracker

* Hadoop Distributed File System: (HDFS)



* MapReduce:



Example : Matrix-Vector Multiplication

Understanding: M is a matrix with dimensions $n \times n$
 v is a vector with n elements

$$\therefore M \times v = n \times 1 \text{ vector}$$

$$\Rightarrow Mv = \sum_{j=1}^n m_{ij} \cdot v_j$$

understand that n is very large ($n \gg 1000$)

↳ Hence cannot be stored on disk

↳ \therefore it is stored into parts onto diff. locations.

Matrix - Vector Multiplication

- Multiply $M = (m_{ij})$ (an $n \times n$ matrix) and $v = (v_i)$ (an n -vector)

$Mv = (x_i)$

$x_i = \sum_{j=1}^n m_{ij} v_j$

Case 1: Large n , M does not fit into main memory, but v does

- Since v fits into main memory, all of v can be copied to every map worker
- Map:** for each matrix element m_{ij} emit key value pair $(i, m_{ij}v_j)$
- Shuffle:** groups all $m_{ij}v_j$ values together for the same i
- Reduce:** sum $m_{ij}v_j$ for all j for the same i

Need all $m_{ij}v_j$ values together for the same i
 So, i needs to be the **key**

Case 1

Matrix - Vector Multiplication

- Multiply $M = (m_{ij})$ (an $n \times n$ matrix) and $v = (v_i)$ (an n -vector)

$Mv = (x_i)$

$x_i = \sum_{j=1}^n m_{ij} v_j$

Case 2: Very large n , even v does not fit into main memory of the nodes

- For every map, many accesses to disk (for parts of v) required!
- Solution:**
 - Partition v so that each partition of v fits into memory
 - Take dot product of one partition of v and the corresponding partition of M
 - Map and reduce same as before, but several times to complete the computation

Case 2

Selection using MapReduce

- SQL: select * from R where C
- Trivial using MapReduce
- Map:** For each tuple t in R , test if t satisfies condition C . If so, produce the key-value pair $(t, 1)$.
- Reduce:** The identity function. It simply passes each key t to the output.

Selection is like "filtering"

Union using MapReduce

- Union of two relations R and S
- Suppose R and S have the same schema
- Map tasks are generated from chunks of both R and S**
- Map:** For each tuple t , produce the key-value pair $(t, 1)$
- Reduce:** Only need to remove duplicates

For all key t , there would be either one (if the tuple is present in one relation) or two values (if the tuple is present in both)

Output t in either case

Reducing is simply emitting 8 keys (after groupby)

R			
$Attr_1$	$Attr_2$	$Attr_3$	$Attr_4$
xyz	abc	1	true
abc	xyz	1	true
xyz	def	1	false
bcd	def	2	true

S			
$Attr_1$	$Attr_2$	$Attr_3$	$Attr_4$
xyz	abc	1	true
abc	xyz	1	false
xyz	def	1	false
bcd	def	2	true

Natural join using MapReduce

- Join $R(A, B)$ with $S(B, C)$ on attribute B
- Map:**
 - For each tuple $t = (a, b)$ of R , emit key value pair $(b, (R, a))$
 - For each tuple $t = (b, c)$ of S , emit key value pair $(b, (S, c))$
- Reduce:**
 - Each key b would be associated with a list of values that are of the form (R, a) or (S, c)
 - Construct all pairs consisting of one with first component R , and the other with first component S , say (R, a) and (S, c) . The output from this key and value list is a sequence of key-value pairs

R		S	
A	B	B	C
x	m	m	1
y	n	p	3
x	p	q	4
w	q	r	7

$R \bowtie S$			
A	B	C	
x	m	1	
x	p	3	
w	q	4	

Grouping and Aggregation using MapReduce

- Group and aggregate on a relation $R(A, B)$ using aggregation function $\gamma(B)$, group by A
- Map:**
 - Grouping comes for free with MapReduce, so the attribute "group by" needs to be the key
 - Map is essentially the identity function
- For each tuple $t = (a, b)$ of R , emit key value pair (a, b)
- Reduce:**
 - Reduce does the aggregation part
 - For the group $\{(a, b_1), \dots, (a, b_k)\}$ represented by a key a , apply γ to obtain $b_2 = \gamma(b_1, \dots, b_k)$


??

R	
A	B
x	2
y	1
z	4
z	1
x	5

select A , sum(B)
 from R group by A

A	sum(B)
x	7
y	1
z	5

Matrix multiplication using MapReduce




- Think of a matrix as a relation with three attributes
- For example matrix A is represented by the relation $A(I, J, V)$
 - For every non-zero entry (i, j, a_{ij}) , the row number is the value of I , column number is the value of J , the entry is the value in V
 - Also advantage: usually most large matrices would be sparse, the relation would have less number of entries
- The product is *analogous* to a natural join followed by a grouping with aggregation

MapReduce

Debagryin Majumdar

Matrix multiplication using MapReduce



- First MapReduce process has produced key value pairs $((i, k), (a_{ij} b_{jk}))$
- Another MapReduce process to group and aggregate
 - Map: identity, just emit the key value pair $((i, k), (a_{ij} b_{jk}))$
 - Reduce:
 - for each key (i, k)
 - produce the sum of the all the values for the key:

MapReduce

Debagryin Majumdar

Combiners


- If the aggregation operator in the reduce algorithm is *associative* and *commutative*
 - Associative: $(a * b) * c = a * (b * c)$
 - Commutative: $(a * b) = (b * a)$
 - It does not matter in which order we aggregate
- We can do some work for reduce in the map step
- Example: the word count problem
 - Instead of emitting $(w, 1)$ for each word
 - We can emit (w, n) for each document or each map task where n is the number of times the word w appears in the whole input chunk

MapReduce

Debagryin Majumdar

Extensions of MapReduce

- MapReduce: **two-step workflow**
- Extension: any collection of functions: acyclic graph representing the workflow



- Each function of the workflow can be executed by many tasks
- A master controller divides the work among the tasks
- Each task's output goes to the successor as input
- Two experimental systems: Clustera (University of Wisconsin), Hydracks (University of California at Irvine)
- Similarly, recursive use of MapReduce

MapReduce

Debagryin Majumdar

